

ProblemStatement: Which model will sutable(bestfit) for the given dataset

In []:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

Data collection

In [2]:

```
1 df=pd.read_csv(r"C:\Users\magam\Downloads\insurance.csv")
2 df.head()
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

Data preprocessing

In [3]:

```
1 df.isnull().sum()
```

Out[3]:

```
age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

In [4]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [5]:

```
1 df.describe()
```

Out[5]:

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

In [6]:

```
1 df['bmi'].value_counts()
```

Out[6]:


```
bmi
32.300    13
28.310     9
30.495     8
30.875     8
31.350     8
..
46.200     1
23.800     1
44.770     1
32.120     1
30.970     1
Name: count, Length: 548, dtype: int64
```

In [62]:

```
1 x=df['bmi']
2 y=df['charges']
3 plt.bar(x,y)
```

Out[62]:

<BarContainer object of 1338 artists>



In [8]:

```
1 df.columns
```

Out[8]:


```
Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'],  
      dtype='object')
```

In [9]:

```
1 sns.barplot(df)
```

Out[9]:

<Axes: >



In [10]:

```
1 df.describe()
```

Out[10]:

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

In [11]:

```
1 df.columns
```

Out[11]:

```
Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'],  
      dtype='object')
```

In [12]:

```
1 sex={"sex":{"male":1,"female":0}}
2 df=df.replace(sex)
3 df
```

Out[12]:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	yes	southwest	16884.92400
1	18	1	33.770	1	no	southeast	1725.55230
2	28	1	33.000	3	no	southeast	4449.46200
3	33	1	22.705	0	no	northwest	21984.47061
4	32	1	28.880	0	no	northwest	3866.85520
...
1333	50	1	30.970	3	no	northwest	10600.54830
1334	18	0	31.920	0	no	northeast	2205.98080
1335	18	0	36.850	0	no	southeast	1629.83350
1336	21	0	25.800	0	no	southwest	2007.94500
1337	61	0	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

In [13]:

```
1 smoker={"smoker":{"yes":1,"no":0}}
2 df=df.replace(smoker)
3 df
```

Out[13]:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	southwest	16884.92400
1	18	1	33.770	1	0	southeast	1725.55230
2	28	1	33.000	3	0	southeast	4449.46200
3	33	1	22.705	0	0	northwest	21984.47061
4	32	1	28.880	0	0	northwest	3866.85520
...
1333	50	1	30.970	3	0	northwest	10600.54830
1334	18	0	31.920	0	0	northeast	2205.98080
1335	18	0	36.850	0	0	southeast	1629.83350
1336	21	0	25.800	0	0	southwest	2007.94500
1337	61	0	29.070	0	1	northwest	29141.36030


1338 rows × 7 columns

In [14]:

```
1 idf=df[['age','sex','bmi','charges']]
2 plt.figure(figsize=(4,4))
3 sns.heatmap(idf.corr(),annot=True)
```

Out[14]:

<Axes: >



In [20]:

```
1 x=df[['age','sex','bmi','children','smoker']]
2 y=df['charges']
```

LINEAR REGRESSION

In [21]:

```
1 from sklearn.linear_model import LinearRegression
2 from sklearn.model_selection import train_test_split
3 lr=LinearRegression()
4 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=100)
5
```


In [22]:

```
1 from sklearn.linear_model import LinearRegression
2 lr=LinearRegression()
3 lr.fit(x_train,y_train)
4 print(lr.intercept_)
5 coeff_df=pd.DataFrame(lr.coef_,x.columns,columns=['coefficient'])
6 coeff_df
```

-10719.483493479494

Out[22]:

	coefficient
age	259.757578
sex	18.216925
bmi	277.903898
children	461.169867
smoker	23981.741027

In [24]:

```
1 score=lr.score(x_test,y_test)
2 print(score)
```

0.780095696440481

In [25]:


```
1 predictions=lr.predict(x_test)
```

In [26]:

```
1 plt.scatter(y_test,predictions)
```

Out[26]:

<matplotlib.collections.PathCollection at 0x1a3c13ada50>



In [27]:

```
1 x=np.array(df['smoker']).reshape(-1,1)
2 y=np.array(df['charges']).reshape(-1,1)
3 df.dropna(inplace=True)
```

In [28]:


```
1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=100)
2 lr.fit(x_train,y_train)
3 lr.fit(x_train,y_train)
```

Out[28]:

```
LinearRegression
LinearRegression()
```

In [29]:

```
1 y_pred=lr.predict(x_test)
2 plt.scatter(x_test,y_test,color='r')
3 plt.plot(x_test,y_pred,color='y')
4 plt.show()
```



LOGISTIC REGRESSION

In [43]:

```
1 from sklearn.linear_model import LogisticRegression
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=100)
3 x=np.array(df['charges']).reshape(-1,1)
4 y=np.array(df['smoker']).reshape(-1,1)
5 df.dropna(inplace=True)
6 lg=LogisticRegression(max_iter=1000)
```

In [44]:

```
1 lg.fit(x_train,y_train)
```

C:\Users\magam\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

Out[44]:

```
▼ LogisticRegression
LogisticRegression(max_iter=1000)
```

In [45]:

```
1 score=lg.score(x_test,y_test)
2 print(score)
```


0.900497512437811

In [46]:

```
1 sns.regplot(x=x,y=y,data=df,logistic=True,ci=None)
```

Out[46]:

<Axes: >



DECISION TREE

In [49]:

```
1 from sklearn.tree import DecisionTreeClassifier
2 clf=DecisionTreeClassifier(random_state=0)
3 clf.fit(x_train,y_train)
```

Out[49]:

```
▼ DecisionTreeClassifier
DecisionTreeClassifier(random_state=0)
```

In [50]:

```
1 clf.fit(x_train,y_train)
```

Out[50]:

```
▼ DecisionTreeClassifier
DecisionTreeClassifier(random_state=0)
```

In [51]:

```
1 score=clf.score(x_test,y_test)
2 print(score)
```

0.8955223880597015

RANDOM FOREST

In [52]:

```
1 from sklearn.ensemble import RandomForestClassifier
2 rfc=RandomForestClassifier()
3 rfc.fit(x_train,y_train)
```

C:\Users\magam\AppData\Local\Temp\ipykernel_920\2210184639.py:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
rfc.fit(x_train,y_train)
```

Out[52]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [53]:

```
1 params={'max_depth':[2,3,5,10,20],
2         'min_samples_leaf':[5,10,20,50,100,200],
3         'n_estimators':[10,25,30,50,100,200]}
```

In [54]:

```
1 from sklearn.model_selection import GridSearchCV
2 grid_search=GridSearchCV(estimator=rfc,param_grid=params,cv=2,scoring="accuracy")
```

In [55]:

```
1 grid_search.fit(x_train,y_train)
```

Please change the shape of y to (n_samples,), for example using `ravel()`.

```
estimator.fit(X_train, y_train, **fit_params)
```

C:\Users\magam\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using `ravel()`.

```
estimator.fit(X_train, y_train, **fit_params)
```

C:\Users\magam\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using `ravel()`.

```
estimator.fit(X_train, y_train, **fit_params)
```

C:\Users\magam\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using `ravel()`.

In [56]:

```
1 grid_search.best_score_
```

Out[56]:

0.9337606837606838

In [57]:

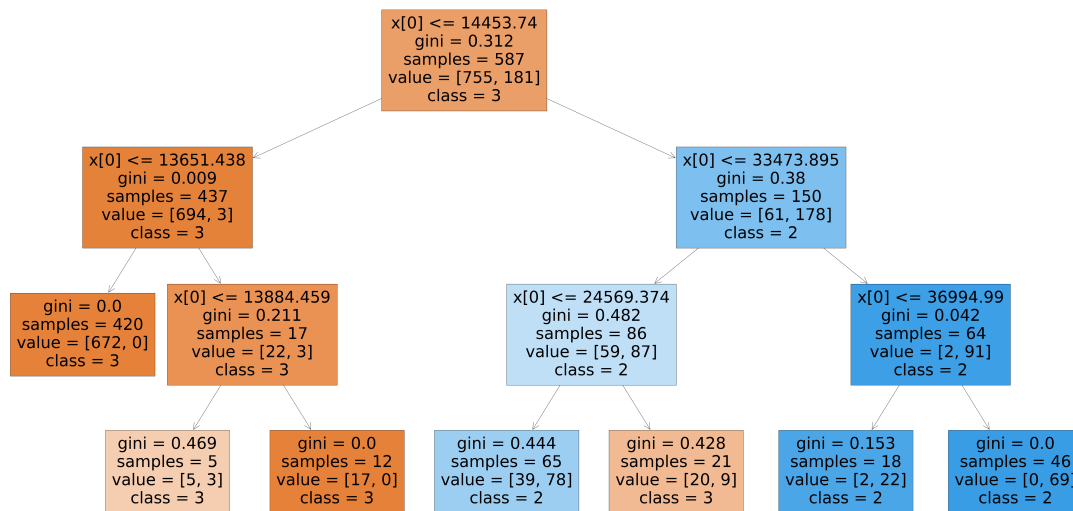
```
1 rf_best=grid_search.best_estimator_  
2 rf_best
```

Out[57]:

```
RandomForestClassifier  
RandomForestClassifier(max_depth=3, min_samples_leaf=5, n_estimators=1  
0)
```

In [58]:

```
1 from sklearn.tree import plot_tree  
2 plt.figure(figsize=(80,40))  
3 plot_tree(rf_best.estimators_[4],class_names=['3','2','1','0'],filled=True);
```



In [59]:

```
1 score=rfc.score(x_test,y_test)  
2 print(score)
```

0.8955223880597015

CONCLUSION

By analysing the data with Linear, logistic, DecisionTree, RandomForest models, I can get 78% for Linear, 90% for Logistic, 89% for DecisionTree and 89% for RandomForest. By this I conclude that Logistical model is the best fit model.

